

Documentation for the **fredpy** Python package

Brian C. Jenkins*
Department of Economics
University of California, Irvine

June 23, 2015

Abstract

fredpy is a Python package for retrieving and working with data from Federal Reserve Economic Data (FRED). The package makes it easy to download specific data series and provides a set of tools for transforming the data in order to construct plots and perform statistical analysis. The **fredpy** package is useful for anyone doing empirical research using the data from FRED and for anyone, e.g. economics teachers, students, and journalists, that will benefit from having an efficient and flexible way to access FRED with Python. **fredpy** is compatible with Python 2.7 and 3.4.

*Email: bcjenkin@uci.edu This project is ongoing and I welcome all feedback. I have no affiliation with the Federal Reserve Bank of St. Louis or with any other component of the Federal Reserve System.

Contents

1 Introduction

Federal Reserve Economic Data or FRED is a rich database maintained by the Federal Reserve Bank of St. Louis.¹ `fredpy` is a Python package that simplifies the process of downloading and manipulating data from FRED. The package offers a streamlined way of retrieving data directly from the FRED website and provides a number of tools to assist with management of the series obtained. This package is particularly useful for creating Python programs that integrate data retrieval with statistical analysis. The package is also well-suited for use with programs that will update figures and tables as new data are released.

2 Preliminaries

Install the `fredpy` package by running `pip install fredpy` in the shell or by obtaining the tarball from [PyPI](#). `fredpy` has the following package dependencies:

- `matplotlib`
- `numpy`
- `scipy`
- `statsmodels`

2.1 Initialization

To create a `fredpy.series` instance, you must have the unique Series ID for the data that you wish to retrieve. Generally, you will find this by searching the FRED website for the data series by name. For example, use the FRED site to find that `GDPC96` is the Series ID for real GDP of the US (quarterly frequency, seasonally adjusted). Create a `fredpy.series` instance based on this data:

```
>>>from fredpy import series
>>>gdp = series('GDPC96')
```

Now the `series` instance `gdp` is available for manipulation.

Each `fredpy.series` instance is initialized with a set of attributes that describe the characteristics of the series retrieved from FRED. These attributes are:

title: Title of the data.

source: Original source of the data.

season: Indicates whether the data has been seasonally adjusted.

freq: Equals 365, 52, 12, 4, or 1 to indicate daily, monthly, quarterly, or annual frequency.

units: Units of the data.

daterange: Date range of data.

updated: Date on which data was last updated.

idCode: The unique FRED identification code for the series.

¹Site: <http://research.stlouisfed.org/fred2/>

data: A list containing the data.

dates: A list containing date strings in yyyy-mm-dd format.

datenumbers: A list containing date numbers to use with the `date time` package.

In terms of the GDP example, find the precise title of data series used by FRED:

```
>>> gdp.title
'Real Gross Domestic Product, 3 Decimal'
```

Or learn that the original source for the data series was the BEA:

```
>>> gdp.source
'U.S. Department of Commerce: Bureau of Economic Analysis'
```

Or find out that quarterly real GDP data goes back to January 1947 and that, as of the time this document was written, the most recent observation was for the quarter beginning April 2014:

```
>>> gdp.daterange
'Range: 1947-01-01 to 2014-04-01'
```

Several of the methods described in the next section will alter the values of the attributes created upon initialization. Furthermore, some methods – the filtering methods in particular – add a couple of new attributes as necessary.

2.2 Methods

Each `fredpy.series` instance has the following methods:

`pc(log=True)`

Replaces the **data** attribute with the percentage change of the data series from the pervious time period. If **log=True**, then percentage change is computed as the difference between the log of the current value and the log of the one period lag value:

$$100 \times [\log(X_t) - \log(X_{t-1})]. \quad (1)$$

If **log=False**, then the percentage change is computed in the standard way:

$$100 \times [(X_t - X_{t-1})/X_{t-1}]. \quad (2)$$

Also, method drops the first elements of the **dates** and **datenumbers** attributes to account for the shorter observation window. Other attributes modified: **units** and **title**.

`apc(log=True)`

Replaces the **data** attribute with the *annual* percentage change of the data series from the pervious time period. If **log=True**, then percentage change is computed as the difference between the log of the current value and the log of the one period lag value:

$$100 \times [\log(X_t) - \log(X_{t-k})], \quad (3)$$

where $k = 1, 4, 12$, or 365 is the annual frequency of the data series. If **log=False**, then the percentage change is computed in the standard way:

$$100 \times [(X_t - X_{t-k})/X_{t-k}]. \quad (4)$$

Also, method drops the first k elements of the `dates` and `datenumbers` attributes to account for the shorter observation window. Other attributes modified: `units` and `title`.

`ma2side(length)`

Constructs a two-sided moving average with window of size $2 \times \text{length} + 1$. *Note the different treatment of `length` relative to the `ma1side` method.* Adds attributes `ma2data`, `ma2dates`, `ma2datenumbers`, and `ma2daterange`.

`ma1side(length)`

Constructs a one-sided (left side) moving average with window of size `length`. *Note the different treatment of `length` relative to the `ma2side` method.* Adds attributes `ma1data`, `ma1dates`, `ma1datenumbers`, and `ma1daterange`.

`recent(N=10)`

Constrains data series to the most recent N observations. The `dates` and `datenumbers` attributes are adjusted from both ends to account for the reduced observation window. Other attributes modified: `daterange`

`window(win)`

Constrains data series to a specified window. `win` is a list specifying the minimum and maximum dates to be included. Dates must be in either mm-dd-yyyy or yyyy-mm-dd format. For example, to restrict observations to those between January 1, 1980 to December 31, 1999:

```
win=['01-01-1980','12-31-1999'].
```

 (5)

Other attributes modified: `daterange`

`log()`

Replaces the `data` attribute with the log of the data series. Other attributes modified: `title`

`bpfilter(low=6,high=32,K=12)`

Computes the bandpass (Baxter-King) filter of the series using the `statsmodels` package. Instead of replacing the modifying the original attributes, this method creates three new attributes:

`bpcycle`: cyclical component of series

`bpdates`: dates corresponding to bp filtered data

`bpdatenumbers`: date numbers corresponding to bp filtered data

Method displays a warning if default values are used for `low`, `high`, and `K` and the original series is not quarterly.

`hpfilter(lamb=1600)`

Computes the Hodrick-Prescott filter of the series using the `statsmodels` package. Instead of replacing the modifying the original `data` attribute, this method creates two new attributes:

hpcycle: cyclical component of series

hptrend: trend component of series

Method displays a warning if default value is used for **lamb** and the original series is not quarterly.

cffilter(low=6,high=32,drift=True)

Computes the Christiano-Fitzgerald filter of the series using the **statsmodels** package. Instead of replacing the modifying the original **data** attribute, this method creates two new attributes:

cfcycle: cyclical component of series

cftrend: trend component of series

Method displays a warning if default values are used for **low**, **high**, and **drift** and the original series is not quarterly.

lntrend()

Computes the linear trend of the series. Instead of replacing the modifying the original **data** attribute, this method creates two new attributes:

lncycle: cyclical component of series

lntrend: trend component of series

Method displays a warning if the original series is not quarterly.

firstdiff()

Computes the first-difference of the series. Instead of replacing the modifying the original **data** attribute, this method creates four new attributes:

diffcycle: cyclical component of series

difftrend: trend component of series

diffdates: shorter date sequence

diffdatenumbers: shorter date numbers

diffdata: shorter data series

monthtoquarter(method='average')

Converts series with monthly frequency to quarterly frequency. Modifies **data**, **dates**, **datenumbers**, **t** attributes. If **method='average'**, then the value for a quarter is the average over each month of the corresponding three month period. If **method='sum'**, then the value for a quarter is the sum over the months of the corresponding three month period. And if **method='end'**, then the value for a quarter is value of the last month of the corresponding three month period.

`quartertoannual(method='average')`

Converts series with quarterly frequency to annual frequency. Modifies `data`, `dates`, `datenumbers`, `t` attributes. If `method='average'`, then the value for a year is the average over each quarter of the corresponding four quarter period. If `method='sum'`, then the value for a year is the sum over the quarters of the corresponding four quarter period. And if `method='end'`, then the value for a year is value of the last quarter of the corresponding four quarter period.

`monthtoannual(method='average')`

Converts series with monthly frequency to annual frequency. Modifies `data`, `dates`, `datenumbers`, `t` attributes. If `method='average'`, then the value for a year is the average over each month of the corresponding twelve month period. If `method='sum'`, then the value for a year is the sum over the months of the corresponding twelve month period. And if `method='end'`, then the value for a year is value of the last month of the corresponding twelve month period.

`percapita(total_pop=True)`

Replaces the `data` attribute with the values divided by the US population using one of two definitions of population:

`total_pop=True`: total population US population.

`total_pop=False`: civilian noninstitutional population is defined as persons 16 years of age and older.

Other attributes modified: `title` and `units`

`recessions()`

Creates gray recession bars for plots. Should be used after a plot has been made but before either (1) a new plot is created or (2) a `show` command is issued.

2.3 Other Functions

The `fredpy` package is equipped with two functions that are occasionally useful for working with `fredpy.series` objects:

`quickplot(x,year_mult=10,show=True,recess=False,save=False,name='file',width=2)`

This function produces a plot of the data for a given `fred` instance `x`. `year_mult` specifies the multiple in which years on the horizontal axis should be formed. `show=True` means that a new plot window should be created. `recess=False` means that no recession bars should be plotted and `recess=True` means the opposite. If `save=True`, the plot is saved in .png format. `width` specifies the width of the plotted line.

`window_equalize(series_list)`

Equalizes the data windows for a list of `fredpy.series` instances by finding the largest date range for which there is data available from each series in `series_list`.

3 Examples


3.1 Example 1: Plotting with the quickplot function

To begin, let's construct a plot of all currently available data for the US unemployment rate. To do this, we run a script that contains the following:

```
from fredpy import series, quickplot
unemp = series('UNRATE')
quickplot(unemp, recess=True, save=True, name='fredpy_example1')
```

This program will display a plot and will save the plot to a file called `fredpy_example1.png` in the working directory. The output is depicted in figure ???. The quick plot function has the advantage of being able to generate plots on the fly, but most of the time you will want to manage plotting directly.

Figure 1: A plot of the US unemployment rate using the `quickplot` function.



3.2 Example 2: Plotting Multiple Series

Next we will construct a more elaborate plot using the `fredpy`. The objective is to create a four-panel plot containing the rate of real GDP growth, CPI inflation, the unemployment rate, and the 3-month T-bill rate. Since the available ranges for these series are different, we will make use of the `window_equalize` function to make sure that we plot all four series over the same date range. First, the import statements:


```
import matplotlib.pyplot as plt
from fredpy import series, window_equalize
import matplotlib.dates as dts
from datetime import dates
```

Next, download data from FRED:

```
gdp = series('GDPC96')
cpi = series('CPIAUCSL')
unemp=series('UNRATE')
tbill=series('TB3MS')
```

Then replace divide gdp data by 1,000 to convert units from billions of dollars to trillions of dollars.

```
gdp.data = gdp.data/1000
```

Use the `apc()` method to compute annual growth in real GDP and the CPI. Note that this method computes percentage change *from one year previous*.

```
gdp.apc()
cpi.apc()
```

Now use the `window_equalize` function to set the date ranges to the smallest range common to all four series.

```
series_list = [gdp,cpi,unemp,tbill]
window_equalize(series_list)
```

Finally, plot the data. Take note of the plot syntax. Use the function `plot_date()` to plot the data of each fredobject against its date numbers.

```
fig = plt.figure()
years10 = dts.YearLocator(10)

# plot real gdp growth
ax1 = fig.add_subplot(221)
ax1.plot_date(gdp.datenumbers,gdp.data,'b-',lw = 3)
ax1.xaxis.set_major_locator(years10)}
ax1.yaxis.set_major_formatter(y_format)
ax1.set_ylabel('Trillions of 2009 $')
fig.autofmt\_xdate()
ax1.grid(True)
gdp.recessions()
ax1.set_title('Real GDP')
```

```

# plot cpi inflation
ax2 = fig.add_subplot(222)
ax2.plot_date(cpi.datenumbers, cpi.data, 'b-', lw = 3)
ax2.xaxis.set_major_locator(years10)
ax2.set_ylabel('Percent')
fig.autofmt_xdate()
ax2.grid(True)
cpi.recessions()
ax2.set_title('CPI Inflation')

# plot unemployment rate
ax3 = fig.add_subplot(223)
ax3.plot_date(unemp.datenumbers, unemp.data, 'b-', lw = 3)
ax3.xaxis.set_major_locator(years10)
ax3.set_ylabel('Percent')
fig.autofmt_xdate()
ax3.grid(True)
unemp.recessions()
ax3.set_title('Unemployment Rate')

# plot 3-mo T-bill rate
ax4 = fig.add_subplot(224)
ax4.plot_date(tbill.datenumbers, tbill.data, 'b-', lw = 3)
ax4.xaxis.set_major_locator(years10)
ax4.set_ylabel('Percent')
fig.autofmt_xdate()
ax4.grid(True)
tbill.recessions()
ax4.set_title('3-mo T-bill Rate')

plt.savefig('fredpy_example2.png', bbox_inches='tight')
plt.show()

```

Figure ?? contains the output of this program.

3.3 Example 3: Filtering

In the final example, we will use the Hodrick-Prescott (HP) and the Bandpass (BP) filters to isolate the business cycle components of real US GDP. Specifically, we will filter log per capita real GDP so that the resulting series may be interpreted as the log-deviations from trend.


```

import matplotlib.pyplot as plt
from fredpy import series, window_equalize
import matplotlib.dates as dts

```

Then download real GDP data from FRED and convert into units of log thousands of dollars per person.

Figure 2: Plots of real GDP growth, CPI inflation, the unemployment rate and the 3-moth T-bill rate.



```
gdp = fred('GDPC96')
gdp.percapita()
gdp.data = gdp.data/1000
gdp.log()
```

Apply the filters. The filter defaults are set for quarterly data.

```
gdp.bpfilter()
gdp.hpfilter()
```

Finally, plot the data.

```

fig = plt.figure()
years10 = dts.YearLocator(10)


ax1 = fig.add_subplot(111)
ax1.plot_date(gdp.datenumbers,gdp.hpcycle,'b-',lw = 2)
ax1.plot_date(gdp.bpdnumbers,gdp.bpcycle,'r--',lw = 2)
ax1.xaxis.set_major_locator(years10)
ax1.set_ylabel('Percent')
fig.autofmt_xdate()
ax1.grid(True)
gdp.recessions()
ax1.legend(['HP','BP'],loc='lower right')

plt.savefig('fredpy_example3.png',bbox_inches='tight')
plt.show()

```

The output of the program is contained in Figure ??.

Figure 3: HP and BP filtered log real GDP per capita.



4 Comments

The examples above emphasize the ease with which the **fredpy** package is used for downloading data and making plots. For another example in this spirit, check out my [animation](#) of the US

Treasury yield curve. Like all other **fredpy**-related materials, the code for this animation is also available from my [Github repository](#) or from my personal [website](#).

The **fredpy** package has uses beyond simple data visualization. A researcher doing statistical analysis using series available from FRED can use the package to create a single Python program – or set of programs – that retrieves data from FRED, performs the analysis, and exports to L^AT_EX tables. Such functionality will reduce the researchers’ time spent managing data files and will allow results to be updated easily as new data become available.